

Introduction to regression trees

MACHINE LEARNING WITH TREE-BASED MODELS IN R



Erin LeDell
Instructor

Train a Regression Tree in R

```
rpart(formula = ____,  
      data = ____,  
      method = ____)
```

formula	is in the format: outcome ~ predictor1+predictor2+etc
data=	specifies the dataframe
method	"class" for classification tree "anova" for regression tree
control=	<i>optional</i> parameters for controlling the tree growth

Train/Validation/Test Split

- training set
- validation set
- test set

Let's practice!

MACHINE LEARNING WITH TREE-BASED MODELS IN R

Performance metrics for regression

MACHINE LEARNING WITH TREE-BASED MODELS IN R



Erin LeDell
Instructor

Common metrics for regression

- Mean Absolute Error (MAE)

$$MAE = \frac{1}{n} \sum |actual - predicted|$$

- Root Mean Square Error (RMSE)

$$RMSE = \sqrt{\frac{1}{n} \sum (actual - predicted)^2}$$

Evaluate a regression tree model

```
pred <- predict(object = model, # model object  
                newdata = test) # test dataset
```

```
library(Metrics)  
  
# Compute the RMSE  
rmse(actual = test$response, # the actual values  
      predicted = pred)      # the predicted values
```

```
2.278249
```

Let's practice!

MACHINE LEARNING WITH TREE-BASED MODELS IN R

What are the hyperparameters for a decision tree?

MACHINE LEARNING WITH TREE-BASED MODELS IN R



Erin LeDell
Instructor

Decision tree hyperparameters

?rpart.control

```
rpart.control {rpart}
```

Control for Rpart Fits

Description

Various parameters that control aspects of the `rpart` fit.

Usage


```
rpart.control(minsplit = 20, minbucket = round(minsplit/3), cp = 0.01,  
              maxcompete = 4, maxsurrogate = 5, usesurrogate = 2, xval = 10,  
              surrogatestyle = 0, maxdepth = 30, ...)
```

Decision tree hyperparameters

- **minsplit**: minimum number of data points required to attempt a split
- **cp**: complexity parameter
- **maxdepth**: depth of a decision tree

Cost-Complexity Parameter (CP)

```
plotcp(grade_model)
```



Cost-Complexity Parameter (CP)

```
print(model$cptable)
```

	CP	nsplit	rel error	xerror	xstd
1	0.06839852	0	1.0000000	1.0080595	0.09215642
2	0.06726713	1	0.9316015	1.0920667	0.09543723
3	0.03462630	2	0.8643344	0.9969520	0.08632297
4	0.02508343	3	0.8297080	0.9291298	0.08571411
5	0.01995676	4	0.8046246	0.9357838	0.08560120
6	0.01817661	5	0.7846679	0.9337462	0.08087153
7	0.01203879	6	0.7664912	0.9092646	0.07982862
8	0.01000000	7	0.7544525	0.9407895	0.08399125

Cost-Complexity Parameter (CP)

```
# Prune the model to optimized cp value  
model_opt <- prune(tree = model, cp = cp_opt)
```

Let's practice!

MACHINE LEARNING WITH TREE-BASED MODELS IN R

Grid Search for model selection

MACHINE LEARNING WITH TREE-BASED MODELS IN R



Erin LeDell
Instructor

Grid Search

- What is a model hyperparameter?
- What is a "grid"?
- What is the goal of a grid search?
- How is the best model chosen?

Set up the grid

```
# Establish a list of possible  
# values for minsplit & maxdepth
```

```
splits <- seq(1, 30, 5)  
depths <- seq(5, 40, 10)
```

```
# Create a data frame containing  
# all combinations
```

```
hyper_grid <- expand.grid(  
  minsplit = splits  
  maxdepth = depths
```

```
hyper_grid[1:10,]
```

	minsplit	maxdepth
1	1	5
2	6	5
3	11	5
4	16	5
5	21	5
6	26	5
7	1	15
8	6	15
9	11	15
10	16	15

Grid Search in R: Train models

```
# Create an empty list to store models
models <- list()
```

```
# Execute the grid search
for (i in 1:nrow(hyper_grid)) {
  # Get minsplitt, maxdepth values at row i
  minsplitt <- hyper_grid$minsplitt[i]
  maxdepth <- hyper_grid$maxdepth[i]

  # Train a model and store in the list
  models[[i]] <- rpart(formula = response ~ .,
                        data = train,
                        method = "anova",
                        minsplitt = minsplitt,
                        maxdepth = maxdepth)
```

```
# Create an empty vector to store RMSE values
rmse_values <- c()
```

```
# Compute validation RMSE
for (i in 1:length(models)) {

  # Retrieve the i^th model from the list
  model <- models[[i]]

  # Generate predictions on grade_valid
  pred <- predict(object = model,
                  newdata = valid)

  # Compute validation RMSE and add to the
  rmse_values[i] <- rmse(actual = valid$response,
                        predicted = pred)

}
```

Let's practice!

MACHINE LEARNING WITH TREE-BASED MODELS IN R