

Hyperparameter tuning in caret

HYPERPARAMETER TUNING IN R



Dr. Shirin Elsinghorst
Senior Data Scientist

Voter dataset from US 2016 election

- Split into training and test set

```
library(tidyverse)
glimpse(voters_train_data)
```

```
Observations: 6,692
Variables: 42
$ turnout16_2016      <chr> "Did not vote", "Did not vote", "Did not vote", "[
$ RIGGED_SYSTEM_1_2016 <int> 2, 2, 3, 2, 2, 3, 3, 1, 2, 3, 4, 4, 4, 3, 1, 2, 2,
$ RIGGED_SYSTEM_2_2016 <int> 3, 3, 2, 2, 3, 3, 2, 2, 1, 2, 4, 2, 3, 2, 3, 4, 3,
$ RIGGED_SYSTEM_3_2016 <int> 1, 1, 3, 1, 1, 1, 2, 1, 1, 2, 1, 2, 1, 2, 1, 1, 1,
$ RIGGED_SYSTEM_4_2016 <int> 2, 1, 2, 2, 2, 2, 2, 2, 1, 3, 3, 1, 3, 3, 1, 3, 3,
$ RIGGED_SYSTEM_5_2016 <int> 1, 2, 2, 2, 2, 3, 1, 1, 2, 3, 2, 2, 1, 3, 1, 1, 2,
$ RIGGED_SYSTEM_6_2016 <int> 1, 1, 2, 1, 2, 2, 2, 1, 2, 2, 1, 3, 1, 3, 1, 1, 1,
$ track_2016          <int> 2, 2, 2, 1, 2, 2, 2, 2, 2, 1, 2, 1, 2, 1, 1, 2, 2,
...
```

Let's train another model with caret

- Stochastic Gradient Boosting

```
library(caret)
library(tictoc)
fitControl <- trainControl(method = "repeatedcv", number = 3, repeats = 5)
tic()
set.seed(42)
gbm_model_voters <- train(turnout16_2016 ~ .,
                          data = voters_train_data,
                          method = "gbm",
                          trControl = fitControl,
                          verbose = FALSE)

toc()
```

32.934 sec elapsed

Let's train another model with caret

```
gbm_model_voters
```

```
Stochastic Gradient Boosting
```

```
...
```

```
Resampling results across tuning parameters:
```

interaction.depth	n.trees	Accuracy	Kappa
1	50	0.9604603	-0.0001774346

```
...
```

```
Tuning parameter 'shrinkage' was held constant at a value of 0.1
```

```
Tuning parameter 'n.minobsinnode' was held constant at a value of 10
```

```
Accuracy was used to select the optimal model using the largest value
```

```
The final values used for the model were n.trees = 50,
```

```
interaction.depth = 1, shrinkage = 0.1 and n.minobsinnode = 10.
```

Cartesian grid search with caret

- Define a **Cartesian grid** of hyperparameters:

```
man_grid <- expand.grid(n.trees = c(100, 200, 250), interaction.depth = c(1, 4),
                        shrinkage = 0.1, n.minobsinnode = 10)

fitControl <- trainControl(method = "repeatedcv", number = 3, repeats = 5)
tic()
set.seed(42)
gbm_model_voters_grid <- train(turnout16_2016 ~ .,
                               data = voters_train_data,
                               method = "gbm",
                               trControl = fitControl,
                               verbose = FALSE,
                               tuneGrid = man_grid)

toc()
```

85.745 sec elapsed

Cartesian grid search with caret

```
gbm_model_voters_grid
```

```
Stochastic Gradient Boosting
```

```
...
```

```
Resampling results across tuning parameters:
```

interaction.depth	n.trees	Accuracy	Kappa
1	100	0.9603108	0.000912769

```
...
```

```
Tuning parameter 'shrinkage' was held constant at a value of 0.1
```

```
Tuning parameter 'n.minobsinnode' was held constant at a value of 10
```

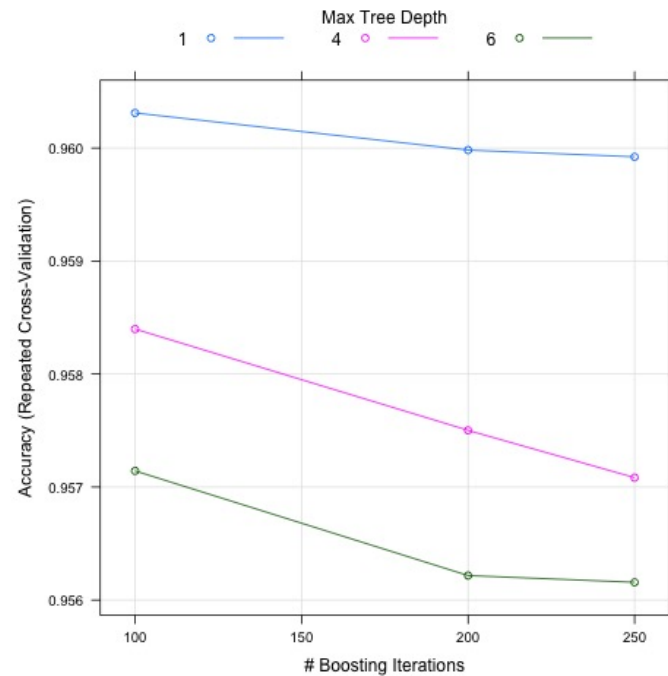
```
Accuracy was used to select the optimal model using the largest value of
```

```
The final values used for the model were n.trees = 100,
```

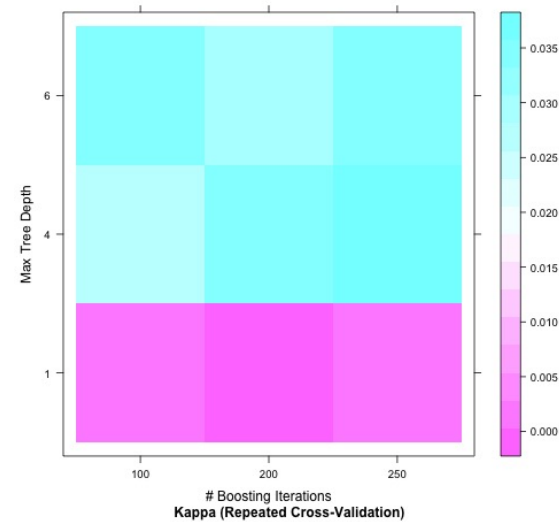
```
interaction.depth = 1, shrinkage = 0.1 and n.minobsinnode = 10.
```

Plot hyperparameter models

```
plot(gbm_model_voters_grid)
```



```
plot(gbm_model_voters_grid,  
     metric = "Kappa",  
     plotType = "level")
```



Test it out for yourself!

HYPERPARAMETER TUNING IN R

Grid vs. Random Search

HYPERPARAMETER TUNING IN R



Dr. Shirin Elsinghorst
Senior Data Scientist

Grid search continued

```
man_grid <- expand.grid(n.trees = c(100, 200, 250), interaction.depth = c(1, 4),
                      shrinkage = 0.1, n.minobsinnode = 10)

fitControl <- trainControl(method = "repeatedcv", number = 3,
                          repeats = 5, search = "grid")

tic()
set.seed(42)
gbm_model_voters_grid <- train(turnout16_2016 ~ .,
                              data = voters_train_data,
                              method = "gbm",
                              trControl = fitControl,
                              verbose= FALSE,
                              tuneGrid = man_grid)

toc()
```

85.745 sec elapsed

Grid search with hyperparameter ranges

```
big_grid <- expand.grid(n.trees = seq(from = 10, to = 300, by = 50),  
                      interaction.depth = seq(from = 1, to = 10,  
                                             length.out = 6),  
                      shrinkage = 0.1,  
                      n.minobsinnode = 10)
```

big_grid

	n.trees	interaction.depth	shrinkage	n.minobsinnode
1	10	1.0	0.1	10
2	60	1.0	0.1	10
3	110	1.0	0.1	10
4	160	1.0	0.1	10
5	210	1.0	0.1	10
6	260	1.0	0.1	10
...				
36	260	10.0	0.1	10

Grid search with many hyperparameter options

```
big_grid <- expand.grid(n.trees = seq(from = 10, to = 300, by = 50),
                      interaction.depth = seq(from = 1, to = 10,
                                             length.out = 6),

                      shrinkage = 0.1,
                      n.minobsinnode = 10)

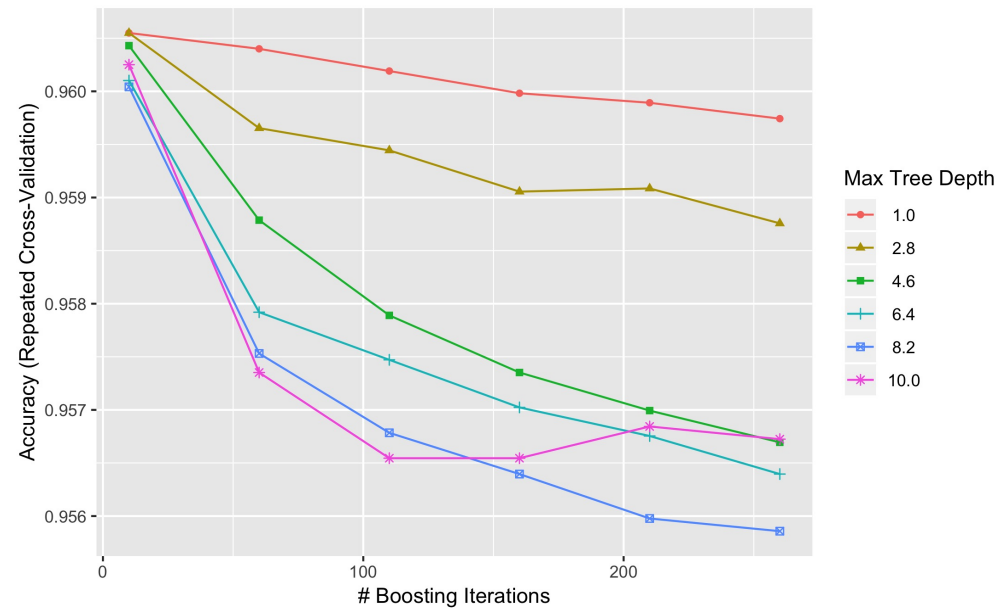
fitControl <- trainControl(method = "repeatedcv", number = 3, repeats = 5, search = "grid")
tic()
set.seed(42)
gbm_model_voters_big_grid <- train(turnout16_2016 ~ .,
                                  data = voters_train_data,
                                  method = "gbm",
                                  trControl = fitControl,
                                  verbose = FALSE,
                                  tuneGrid = big_grid)

toc()
```

240.698 sec elapsed

Cartesian grid vs random search

```
ggplot(gbm_model_voters_big_
```



- Grid search can get **slow** and **computationally expensive** very quickly!
- Therefore, in reality, we often use **random search**.

Random search in caret

```
# Define random search in trainControl function
library(caret)
fitControl <- trainControl(method = "repeatedcv", number = 3, repeats = 5, search = "random")
```

```
# Set tuneLength argument
tic()
set.seed(42)
gbm_model_voters_random <- train(turnout16_2016 ~ .,
                                data = voters_train_data,
                                method = "gbm",
                                trControl = fitControl,
                                verbose = FALSE,
                                tuneLength = 5)
toc()
```

46.432 sec elapsed

Random search in caret

```
gbm_model_voters_random
```

```
Stochastic Gradient Boosting
```

```
...
```

```
Resampling results across tuning parameters:
```

shrinkage	interaction.depth	n.minobsinnode	n.trees	Accuracy	Kappa
0.08841129	4	6	4396	0.9670737	-0.0085331
0.09255042	2	7	540	0.9630635	-0.0132916
0.14484962	3	21	3154	0.9570179	-0.0139702
0.34935098	10	10	2566	0.9610734	-0.0157268
0.43341085	1	13	2094	0.9460727	-0.0247910

```
Accuracy was used to select the optimal model using the largest value.
```

```
The final values used for the model were n.trees = 4396,  
interaction.depth = 4, shrinkage = 0.08841129 and n.minobsinnode = 6.
```

- *Beware:* in `caret` random search can **NOT** be combined with grid search!

Let's get coding!

HYPERPARAMETER TUNING IN R

Adaptive resampling

HYPERPARAMETER TUNING IN R



Dr. Shirin Elsinghorst

Senior Data Scientist

What is Adaptive Resampling?

Grid Search

- **All** hyperparameter combinations are computed.

Random Search

- Random **subsets** of hyperparameter combinations are computed.

→ Evaluation of best combination is done **at the end**.

Adaptive Resampling

- Hyperparameter combinations are **resampled** with values near combinations that performed well.
- Adaptive Resampling is, therefore, **faster and more efficient!**

"Futility Analysis in the Cross-Validation of Machine Learning Models." Max Kuhn;

Adaptive resampling in caret

```
trainControl : method = "adaptive_cv" + search = "random" +  
adaptive =
```

- *min*: minimum number of resamples per hyperparameter
- *alpha*: confidence level for removing hyperparameters
- *method*: "gls" for linear model or "BT" for Bradley-Terry
- *complete*: if TRUE generates full resampling set

```
fitControl <- trainControl(method = "adaptive_cv",  
                           adaptive = list(min = 2, alpha  
                                           method = "gls"  
                                           search = "random"))
```

- `trainControl()` + `tuneLength = x`

```
fitControl <- trainControl(method = "adaptive_cv", number = 3, repeats = 3,  
                           adaptive = list(min = 2,  
                                           alpha = 0.05,  
                                           method = "gls",  
                                           complete = TRUE),  
                           search = "random")  
  
tic()  
set.seed(42)  
gbm_model_voters_adaptive <- train(turnout16_2016 ~ .,  
                                   data = voters_train_data,  
                                   method = "gbm",  
                                   trControl = fitControl,  
                                   verbose = FALSE,  
                                   tuneLength = 7)  
  
toc()
```

1239.837 sec elapsed

Adaptive resampling

```
gbm_model_voters_adaptive
```

```
...
```

```
Resampling results across tuning parameters:
```

shrinkage	interaction.depth	n.minobsinnode	n.trees	Accuracy	Kappa
0.07137493	5	6	4152	0.9564654	0.02856571
0.08408739	5	14	674	0.9547185	0.02098853
0.28552325	8	15	3209	0.9568141	0.03024238
0.33663932	10	13	2595	0.9571130	0.04250979
0.54251480	3	24	3683	0.9482171	0.03568586
0.56406870	7	25	4685	0.9549898	0.05284333
0.58695763	8	24	1431	0.9520286	0.02742592

```
Accuracy was used to select the optimal model using the largest value.
```

```
The final values used for the model were n.trees = 2595,  
interaction.depth = 10, shrinkage = 0.3366393 and n.minobsinnode = 13.
```

Let's get coding!

HYPERPARAMETER TUNING IN R